

Cursor Training Program — Post-Training Assessment (MCQ-Based)

For Working Professionals — After Completing Training

Instructions

Aspect	Details
Time Limit	15-20 minutes
Format	Multiple Choice Questions (MCQ) only
Number of Questions	40 questions
Passing Score	70% (28/40 correct)
Purpose	Measure knowledge gain after completing the training

Name (Optional): _____

Date: _____

Training Cohort: _____

Section 1: Mental Models & Foundations (5 questions)

Q1. You ask an AI model the same question twice and get slightly different answers. Why does this happen?

- A) The AI is malfunctioning
 - B) AI models are probabilistic, not deterministic – they predict likely tokens, not exact answers
 - C) Your internet connection is unstable
 - D) The AI is intentionally being unpredictable
-

Q2. An AI confidently suggests using a function `i2c_write_register_stm32_hal_v3()` that doesn't exist in your codebase. What is this phenomenon called?

- A) Compilation error
 - B) Token limit issue
 - C) Hallucination
 - D) Optimization
-

Q3. Your team wants to reduce AI costs. Which strategy would be most effective?

- A) Always use the most expensive model for better quality
 - B) Use cheaper models for simple tasks and reserve expensive models for complex work
 - C) Increase context window to maximum for every request
 - D) Disable all caching
-

Q4. What is the single most valuable skill when working with AI coding assistants?

- A) Knowing multiple programming languages
 - B) Managing what the model "sees" (context engineering)
 - C) Writing the shortest possible prompts
 - D) Always using the latest model
-

Q5. What is the key difference between a "chatbot" and an "agent"?

- A) Chatbots are free; agents cost money
 - B) Chatbots respond to prompts; agents pursue multi-step goals and take actions
 - C) Chatbots work offline; agents require internet
 - D) There is no meaningful difference
-

Section 2: Cursor Editor Essentials (5 questions)

Q6. You open a large, unfamiliar firmware codebase. What is the most effective first step?

- A) Immediately ask the agent to refactor all drivers
 - B) Ask: "Explain this codebase – main entry points, key modules"
 - C) Manually read every file line by line
 - D) Create a pull request with documentation
-

Q7. You need to add a complex feature across multiple files and want the AI to plan first. What should you use?

- A) Ask Mode
 - B) Plan Mode (Shift+Tab)
 - C) Direct coding request
 - D) Cloud Agent only
-

Q8. What is the purpose of @mentions (e.g., @timers.c @TIM2_IRQHandler) in Cursor?

- A) Tag a team member for review
 - B) Point the agent to exact files, symbols, or context
 - C) Create a hyperlink in comments
 - D) Format code automatically
-

Q9. Before letting an AI make significant changes to your code, what should you create as a safety net?

- A) A new branch in git
 - B) A checkpoint in Cursor
 - C) A backup of the entire project
 - D) All of the above
-

Q10. You ask the AI to change a variable name. The AI also renames other variables and reformats the file. What should you do?

- A) Accept all changes
 - B) Review the diff, reject unwanted changes, and ask the AI to be more focused
 - C) Close the session and start over
 - D) Disable the AI
-

Section 3: Agent Modes & Tools (5 questions)

Q11. What is the difference between Ask Mode and Agent Mode?

- A) Ask Mode is faster; Agent Mode is slower
 - B) Ask Mode is read-only; Agent Mode can make changes and run tools
 - C) Ask Mode uses free models; Agent Mode uses paid models
 - D) There is no difference
-

Q12. Your firmware is crashing. You want the AI to see debug output and diagnose the issue. Which tool should the AI use?

- A) Browser Tool
 - B) Terminal Tool
 - C) Files Tool only
 - D) Cannot do this
-

Q13. You want the AI to inspect a live web dashboard showing sensor data. Which tool should it use?

- A) Terminal Tool
 - B) Browser Tool
 - C) Files Tool
 - D) MCP only
-

Q14. You ask the AI to run `make` and fix compilation errors. The AI requests approval. Why?

- A) The AI is malfunctioning
 - B) Terminal commands require approval by default for safety
 - C) Your API key is invalid
 - D) The command is malformed
-

Q15. Which is an example of an effective, constrained prompt?

- A) "Fix the bug"
 - B) "Fix the IndexError when input list is empty. Do not change function signature. Show the diff."
 - C) "Make the code better"
 - D) "Rewrite everything"
-

Section 4: Customizing Cursor (4 questions)

Q16. Where should you encode team coding standards so the AI follows them automatically?

- A) In a Slack message
 - B) In a `.cursor/rules/*.mdc` file
 - C) In a README.md
 - D) In an email to the team
-

Q17. What is a "Skill" in Cursor?

- A) A keyboard shortcut
 - B) A reusable, specialized workflow invoked by name (e.g., `/pr-review`)
 - C) A type of AI model
 - D) A debugging tool
-

Q18. What does MCP (Model Context Protocol) enable?

- A) Faster code completion
 - B) Connecting external tools (GitHub, Slack, Jira) for the agent to use
 - C) Voice control of Cursor
 - D) Automatic code formatting
-

Q19. What is a Subagent?

- A) A less intelligent AI model
 - B) A separate agent instance that handles specialized tasks in parallel
 - C) A bug in the agent
 - D) A type of prompt template
-

Section 5: CLI & Automation (4 questions)

Q20. How do you run a one-shot AI command from the terminal without entering an interactive session?

- A) `agent`
 - B) `agent "your prompt here"`
 - C) `cursor --gui`
 - D) `ai run`
-

Q21. You need to run a long code analysis task (30+ minutes) and close your laptop. What should you use?

- A) Local Agent in Ask Mode
 - B) Cloud Agent
 - C) Terminal one-shot command
 - D) Cannot run long tasks offline
-

Q22. What does the `/resume` command do in the Cursor CLI?

- A) Restart the current session
 - B) List previous sessions and resume one
 - C) Save the current conversation
 - D) Clear the context window
-

Q23. How do you hand off a local session to a Cloud Agent from the CLI?

- A) Prefix message with `&` (e.g., `& "Run this task"`)
 - B) Use `--cloud` flag
 - C) Use `/cloud` command
 - D) Cannot do this from CLI
-

Section 6: Cloud Agents API & Webhooks (6 questions)

Q24. What are the two required fields when creating a Cloud Agent via API?

- A) `agent.name` and `agent.description`
 - B) `prompt.text` and `repos[].url`
 - C) `model.id` and `temperature`
 - D) `webhookUrl` and `webhookSecret`
-

Q25. What are Server-Sent Events (SSE) used for in the Cloud Agents API?

- A) Encrypting API requests
 - B) Streaming real-time agent responses (thoughts, tool calls, output)
 - C) Batch processing multiple agents
 - D) Scheduled agent execution
-

Q26. How do you verify a webhook genuinely came from Cursor?

- A) Check the sender email address
 - B) Verify the HMAC-SHA256 signature in `X-Webhook-Signature` header
 - C) Trust it because it came over HTTPS
 - D) Compare payload with known example
-

Q27. What is the purpose of ngrok when testing webhooks locally?

- A) Debug API responses
 - B) Expose local webhook server to the internet for Cursor to reach
 - C) Encrypt webhook payloads
 - D) Cache webhook responses
-

Q28. What are "artifacts" in Cloud Agents?

- A) Bugs in the agent's code
 - B) Files produced by the agent (logs, screenshots, generated code)
 - C) The agent's configuration files
 - D) Unused code
-

Q29. How long are artifact download URLs valid?

- A) 1 hour
 - B) 15 minutes
 - C) 24 hours
 - D) 7 days
-

Section 7: Admin & Analytics APIs (5 questions)

Q30. What is the primary difference between a User API Key and an Admin API Key?

- A) User keys expire faster
 - B) Admin keys can access team usage data and manage members
 - C) User keys are free; Admin keys cost extra
 - D) There is no difference
-

Q31. Your manager wants a report showing total AI spending per day for the last 30 days. Which endpoint should you use?

- A) `GET /v1/admin/members`
 - B) `GET /v1/admin/analytics/usage/daily`
 - C) `POST /v1/agents`
 - D) `GET /v1/models`
-

Q32. How do you prevent a team member from exceeding \$50 in AI usage per month?

- A) Configure a webhook alert
 - B) Set a user spend limit via `PATCH /v1/admin/policies/users/{userId}/limits`
 - C) Create a repository instruction
 - D) Create a skill
-

Q33. What is the recommended safe pattern for removing a team member?

- A) Immediately hard delete the account
 - B) Audit → deactivate (soft delete) → transfer ownership → (optional) hard delete
 - C) Just remove their email from the team list
 - D) Revoke only their API keys
-

Q34. When creating a leaderboard for team usage, what is a responsible practice?

- A) Show full names and exact costs
 - B) Anonymize users and focus on positive metrics (efficiency, savings)
 - C) Only show it to managers
 - D) Publish it publicly
-

Section 8: AI Code Tracking & Reporting (3 questions)

Q35. What does the AI commit metrics endpoint measure?

- A) Number of commits per day
 - B) Per-commit attribution of AI-generated vs. human-written lines
 - C) Which AI model was used most frequently
 - D) Speed of code reviews
-

Q36. Your team uses Tableau for dashboards. What is the most efficient way to export AI usage data?

- A) Copy-paste from Cursor UI
 - B) Use streaming CSV export (</admin/analytics/export/csv>)
 - C) Take screenshots
 - D) Manually type data into Excel
-

Q37. What information does a granular AI change event include?

- A) Only the file name
 - B) File path, line range, model, user, timestamp, original and new code
 - C) Only the user who made the change
 - D) Only the timestamp
-

Section 9: Practical Scenarios (3 questions)

Q38. You inherit a legacy firmware codebase with no documentation. What is your first step?

- A) Ask AI to rewrite everything
- B) Ask AI: "Trace the data flow from sensor to application. Identify all files involved."
- C) Manually read every file
- D) Give up and recommend replacing the system

Q39. You need to add a watchdog timer across 5 files. What workflow ensures safety?

- A) Tell AI "Add watchdog timer" and accept all changes
- B) Use Plan Mode → review plan → execute step by step → review each diff
- C) Make all changes manually
- D) Ask a junior engineer to do it

Q40. An auditor asks for a record of all AI-generated code changes in the last 90 days. How do you provide this?

- A) Say you don't have it
- B) Export granular events via `/admin/analytics/events` to CSV
- C) Ask each developer to self-report
- D) Provide only the total cost

Answer Sheet (For Participants to Fill)

Q#	Answer	Q#	Answer	Q#	Answer	Q#	Answer
1		11		21		31	
2		12		22		32	
3		13		23		33	
4		14		24		34	
5		15		25		35	
6		16		26		36	
7		17		27		37	
8		18		28		38	
9		19		29		39	
10		20		30		40	

Answer Key (For Instructors Only)

Q#	Answer	Q#	Answer	Q#	Answer	Q#	Answer
1	B	11	B	21	B	31	B
2	C	12	B	22	B	32	B
3	B	13	B	23	A	33	B
4	B	14	B	24	B	34	B
5	B	15	B	25	B	35	B

Q#	Answer	Q#	Answer	Q#	Answer	Q#	Answer
6	B	16	B	26	B	36	B
7	B	17	B	27	B	37	B
8	B	18	B	28	B	38	B
9	D	19	B	29	B	39	B
10	B	20	B	30	B	40	B

Scoring Summary

Section	Questions	Correct	Score
Section 1: Mental Models	1-5	___ / 5	
Section 2: Editor Essentials	6-10	___ / 5	
Section 3: Agent Modes & Tools	11-15	___ / 5	
Section 4: Customizing Cursor	16-19	___ / 4	
Section 5: CLI & Automation	20-23	___ / 4	
Section 6: Cloud Agents API	24-29	___ / 6	
Section 7: Admin & Analytics	30-34	___ / 5	
Section 8: AI Code Tracking	35-37	___ / 3	
Section 9: Practical Scenarios	38-40	___ / 3	
Total	1-40	___ / 40	

Score Interpretation

Score	Percentage	Interpretation
36-40	90-100%	Excellent
32-35	80-89%	Very Good
28-31	70-79%	Good (Passing)
24-27	60-69%	Needs Review
Below 24	Below 60%	Recommend Retraining

End of Post-Training Assessment (MCQ-Based)